

Assignment_6.1

July 15, 2021

```
[1]: from keras import layers
from keras import models
import pandas as pd

#initiate a small convnet

model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))

#add a clasifier on top of the convnet
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 5, 5, 64)	0	
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928

```

dense_1 (Dense)           (None, 10)          650
=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
-----
```

```
[2]: #training the convnet on MNIST images
from keras.datasets import mnist
from keras.utils import to_categorical
import numpy as np

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

#shuffle the training set
for _ in range(5):
    indexes = np.random.permutation(len(train_images))

train_images = train_images[indexes]
train_labels = train_labels[indexes]

#set aside 10,000 for validation
val_images = train_images[:10000,:]
val_labels = train_labels[:10000,:]

# leave rest in training set
train_images2 = train_images[10000:,:]
train_labels2 = train_labels[10000:,:]

train_images2.shape, val_images.shape
```

```
[2]: ((50000, 28, 28, 1), (10000, 28, 28, 1))
```

```
[3]: model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history = model.fit(train_images2, train_labels2, epochs=5, batch_size=64,
                     validation_data=(val_images, val_labels))
```

```
Epoch 1/5
782/782 [=====] - 13s 16ms/step - loss: 0.4656 -
accuracy: 0.8508 - val_loss: 0.0915 - val_accuracy: 0.9726
Epoch 2/5
782/782 [=====] - 11s 14ms/step - loss: 0.0621 -
accuracy: 0.9819 - val_loss: 0.0573 - val_accuracy: 0.9833
Epoch 3/5
782/782 [=====] - 11s 15ms/step - loss: 0.0349 -
accuracy: 0.9897 - val_loss: 0.0426 - val_accuracy: 0.9869
Epoch 4/5
782/782 [=====] - 11s 14ms/step - loss: 0.0259 -
accuracy: 0.9924 - val_loss: 0.0439 - val_accuracy: 0.9880
Epoch 5/5
782/782 [=====] - 11s 14ms/step - loss: 0.0200 -
accuracy: 0.9935 - val_loss: 0.0444 - val_accuracy: 0.9887
```

```
[4]: history.history.keys()
```

```
[4]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```


```
[5]: import matplotlib.pyplot as plt

train_loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(history.history['loss']) + 1)

plt.plot(epochs, train_loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and Validation Losses')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
plt.savefig('results/6_1_lossplot.png')
```




<Figure size 432x288 with 0 Axes>

```
[6]: train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

epochs = range(1, len(history.history['accuracy']) + 1)

plt.plot(epochs, train_acc, 'bo', label='Training Accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
plt.savefig('results/6_1_accplot.png')
```



<Figure size 432x288 with 0 Axes>

```
[7]: #retrain and evaluate for 3 epochs
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(train_images, train_labels, epochs=3, batch_size=64)
results = model.evaluate(test_images, test_labels)
```

```
Epoch 1/3
938/938 [=====] - 12s 12ms/step - loss: 0.0223 -
accuracy: 0.9938
Epoch 2/3
938/938 [=====] - 12s 12ms/step - loss: 0.0156 -
accuracy: 0.9954
Epoch 3/3
938/938 [=====] - 12s 12ms/step - loss: 0.0116 -
accuracy: 0.9961
313/313 [=====] - 1s 3ms/step - loss: 0.0332 -
accuracy: 0.9908
```

```
[8]: results
```

```
[8]: [0.03317027539014816, 0.9908000230789185]
```

```
[9]: history.history
```

```
[9]: {'loss': [0.02030199207365513, 0.016139207407832146, 0.013412469066679478],  
      'accuracy': [0.9940833449363708, 0.9952166676521301, 0.9956833124160767]}
```

```
[10]: model.save('results/6_1_model.h5')
```

```
[11]: prediction_results = model.predict(test_images)
```

```
[12]: #write metrics to file  
with open('results/6_1_metrics.txt', 'w') as f:  
    f.write('Training Loss: {}'.format(str(history.history['loss'])))  
    f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))  
    f.write('\nTest Loss: {}'.format(results[0]))  
    f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[13]: predictions = pd.DataFrame(prediction_results,  
                                columns=['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])  
predictions.to_csv('results/6_1_predictions.csv', index=False)
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```