

Assignment_6.2b

July 15, 2021

```
[1]: from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
import pandas as pd
import matplotlib.pyplot as plt

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

A local file was found, but it seems to be incomplete or outdated because the auto file hash does not match the original value of
6d958be074577803d12ecdefd02955f39262c83c16fe9348329d7fe0b5c001ce so we will re-download the data.

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 22s 0us/step

```
[2]: x_train.shape, y_train.shape
```

```
[2]: ((50000, 32, 32, 3), (50000, 1))
```

```
[3]: x_test.shape, y_test.shape
```

```
[3]: ((10000, 32, 32, 3), (10000, 1))
```

```
[4]: # Preprocess the data (these are NumPy arrays)
x_train = x_train.astype("float32")
x_test = x_test.astype("float32")

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
[5]: # Reserve 10,000 samples for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train_2 = x_train[:-10000]
y_train_2 = y_train[:-10000]
```

```
[6]: train_datagen = ImageDataGenerator(rescale=1./255,
                                       rotation_range=40,
                                       width_shift_range=0.2,
                                       height_shift_range=0.2,
                                       shear_range=0.2,
                                       zoom_range=0.2,
                                       horizontal_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow(x_train_2, y_train_2, batch_size=32)

validation_generator = train_datagen.flow(x_val, y_val, batch_size=32)
```

```
[7]: #instantiate the model
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D(2,2))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0

```

-----  

flatten (Flatten)           (None, 256)          0  

-----  

dropout (Dropout)           (None, 256)          0  

-----  

dense (Dense)               (None, 64)           16448  

-----  

dense_1 (Dense)              (None, 10)           650  

=====  

Total params: 73,418  

Trainable params: 73,418  

Non-trainable params: 0
-----
```

[8]: `from keras import optimizers`

```
model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

[9]: `history = model.fit_generator(train_generator,
 steps_per_epoch=len(x_train_2) / 32,
 epochs=30,
 validation_data=validation_generator,
 validation_steps=len(x_val) / 32)`


```
/opt/conda/lib/python3.8/site-
packages/tensorflow/python/keras/engine/training.py:1844: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
    warnings.warn(``Model.fit_generator`` is deprecated and `

Epoch 1/30
1250/1250 [=====] - 60s 47ms/step - loss: 2.2345 -
accuracy: 0.1506 - val_loss: 1.9912 - val_accuracy: 0.2555
Epoch 2/30
1250/1250 [=====] - 57s 46ms/step - loss: 1.9917 -
accuracy: 0.2482 - val_loss: 1.9086 - val_accuracy: 0.2968
Epoch 3/30
1250/1250 [=====] - 58s 46ms/step - loss: 1.9217 -
accuracy: 0.2808 - val_loss: 1.8425 - val_accuracy: 0.3159
Epoch 4/30
1250/1250 [=====] - 58s 46ms/step - loss: 1.8709 -
accuracy: 0.3068 - val_loss: 1.8045 - val_accuracy: 0.3331
Epoch 5/30
1250/1250 [=====] - 57s 46ms/step - loss: 1.8241 -
accuracy: 0.3224 - val_loss: 1.7681 - val_accuracy: 0.3492
Epoch 6/30
```

```
1250/1250 [=====] - 58s 46ms/step - loss: 1.7955 -  
accuracy: 0.3340 - val_loss: 1.7300 - val_accuracy: 0.3601  
Epoch 7/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.7610 -  
accuracy: 0.3522 - val_loss: 1.7076 - val_accuracy: 0.3737  
Epoch 8/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.7422 -  
accuracy: 0.3605 - val_loss: 1.7180 - val_accuracy: 0.3708  
Epoch 9/30  
1250/1250 [=====] - 57s 46ms/step - loss: 1.7113 -  
accuracy: 0.3688 - val_loss: 1.6650 - val_accuracy: 0.3983  
Epoch 10/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.7001 -  
accuracy: 0.3827 - val_loss: 1.6621 - val_accuracy: 0.3984  
Epoch 11/30  
1250/1250 [=====] - 57s 46ms/step - loss: 1.6806 -  
accuracy: 0.3853 - val_loss: 1.6346 - val_accuracy: 0.4102  
Epoch 12/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.6613 -  
accuracy: 0.3946 - val_loss: 1.6068 - val_accuracy: 0.4207  
Epoch 13/30  
1250/1250 [=====] - 57s 46ms/step - loss: 1.6567 -  
accuracy: 0.4003 - val_loss: 1.5812 - val_accuracy: 0.4277  
Epoch 14/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.6438 -  
accuracy: 0.4026 - val_loss: 1.5715 - val_accuracy: 0.4371  
Epoch 15/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.6283 -  
accuracy: 0.4103 - val_loss: 1.5709 - val_accuracy: 0.4381  
Epoch 16/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.6083 -  
accuracy: 0.4170 - val_loss: 1.5413 - val_accuracy: 0.4493  
Epoch 17/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.6014 -  
accuracy: 0.4205 - val_loss: 1.5466 - val_accuracy: 0.4400  
Epoch 18/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.5938 -  
accuracy: 0.4247 - val_loss: 1.5225 - val_accuracy: 0.4587  
Epoch 19/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.5859 -  
accuracy: 0.4317 - val_loss: 1.5200 - val_accuracy: 0.4547  
Epoch 20/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.5902 -  
accuracy: 0.4300 - val_loss: 1.5099 - val_accuracy: 0.4598  
Epoch 21/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.5650 -  
accuracy: 0.4408 - val_loss: 1.4952 - val_accuracy: 0.4705  
Epoch 22/30
```

```
1250/1250 [=====] - 58s 46ms/step - loss: 1.5661 -  
accuracy: 0.4425 - val_loss: 1.4884 - val_accuracy: 0.4624  
Epoch 23/30  
1250/1250 [=====] - 58s 46ms/step - loss: 1.5450 -  
accuracy: 0.4460 - val_loss: 1.5389 - val_accuracy: 0.4531  
Epoch 24/30  
1250/1250 [=====] - 57s 45ms/step - loss: 1.5398 -  
accuracy: 0.4496 - val_loss: 1.4618 - val_accuracy: 0.4777  
Epoch 25/30  
1250/1250 [=====] - 57s 46ms/step - loss: 1.5385 -  
accuracy: 0.4493 - val_loss: 1.4465 - val_accuracy: 0.4853  
Epoch 26/30  
1250/1250 [=====] - 57s 46ms/step - loss: 1.5231 -  
accuracy: 0.4564 - val_loss: 1.4627 - val_accuracy: 0.4754  
Epoch 27/30  
1250/1250 [=====] - 53s 43ms/step - loss: 1.5106 -  
accuracy: 0.4621 - val_loss: 1.4703 - val_accuracy: 0.4792  
Epoch 28/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.5147 -  
accuracy: 0.4562 - val_loss: 1.4331 - val_accuracy: 0.4902  
Epoch 29/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.5085 -  
accuracy: 0.4626 - val_loss: 1.4535 - val_accuracy: 0.4820  
Epoch 30/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.4960 -  
accuracy: 0.4672 - val_loss: 1.4255 - val_accuracy: 0.4918
```

```
[10]: train_loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs = range(1, len(history.history['loss']) + 1)  
  
plt.plot(epochs, train_loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and Validation Losses')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
  
plt.show()  
plt.savefig('results/6_2b_lossplot.png')
```




<Figure size 432x288 with 0 Axes>

```
[11]: train_loss = history.history['accuracy']
val_loss = history.history['val_accuracy']

epochs = range(1, len(history.history['accuracy']) + 1)

plt.plot(epochs, train_loss, 'bo', label='Training accuracy')
plt.plot(epochs, val_loss, 'b', label='Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
plt.savefig('results/6_2b_accplot.png')
```



<Figure size 432x288 with 0 Axes>

```
[12]: #retrain the model and evaluate on test
train_generator = train_datagen.flow(x_train, y_train, batch_size=32)

model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

#16 epochs chosen based on graphs above
history = model.fit_generator(train_generator,
                               steps_per_epoch=len(x_train) / 32,
                               epochs=16)
results = model.evaluate(x_test, y_test)
```

```
Epoch 1/16
1562/1562 [=====] - 45s 28ms/step - loss: 1.4980 -
accuracy: 0.4654
Epoch 2/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4841 -
accuracy: 0.4684
Epoch 3/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4770 -
accuracy: 0.4744
```

```
Epoch 4/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4596 -
accuracy: 0.4796
Epoch 5/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4613 -
accuracy: 0.4803
Epoch 6/16
1562/1562 [=====] - 45s 28ms/step - loss: 1.4551 -
accuracy: 0.4803
Epoch 7/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4536 -
accuracy: 0.4842
Epoch 8/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4387 -
accuracy: 0.4878
Epoch 9/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4443 -
accuracy: 0.4859
Epoch 10/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4298 -
accuracy: 0.4911
Epoch 11/16
1562/1562 [=====] - 45s 29ms/step - loss: 1.4355 -
accuracy: 0.4876
Epoch 12/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4225 -
accuracy: 0.4950
Epoch 13/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4078 -
accuracy: 0.4986
Epoch 14/16
1562/1562 [=====] - 45s 29ms/step - loss: 1.4052 -
accuracy: 0.5009
Epoch 15/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4082 -
accuracy: 0.4969
Epoch 16/16
1562/1562 [=====] - 45s 28ms/step - loss: 1.4050 -
accuracy: 0.5023
313/313 [=====] - 1s 4ms/step - loss: 257.6946 -
accuracy: 0.3655
```

```
[13]: model.save('results/6_2b_model.h5')
```

```
[14]: prediction_results = model.predict(x_test)
```

```
[15]: #write metrics to file
with open('results/6_2b_metrics.txt', 'w') as f:
    f.write('Training Loss: {}'.format(str(history.history['loss'])))
    f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
    f.write('\nTest Loss: {}'.format(results[0]))
    f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[16]: predictions = pd.DataFrame(prediction_results,
→columns=['0','1','2','3','4','5','6','7','8','9'])
predictions.to_csv('results/6_2b_predictions.csv', index=False)
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```